



UNIVERSITÉ GRENOBLE ALPES / ENSIMAG

Internship Report

Phage-Host Classification Using Conformal Prediction????

Presented by:

Nizalia SIDDIQUI

Supervisor : Clovis Galiez

June 2026

Abstract

Acknowledgements

ENGLISH VERSION

Version française

Contents

Acknowledgements	3
1 Introduction	1
2 Background & State of the Art	3
2.1 Bacteriophage Host Prediction Problem	3
2.2 State of the Art	3
2.3 Conformal Prediction	3
2.3.1 Setup and nonconformity scores	3
2.3.2 prediction set and coverage guarantee	3
2.3.3 Split conformal prediction	3
2.4 The CSA Method	4
2.4.1 Motivation	4
2.4.2 Two-stage calibration	4
2.4.3 The envelope concept	4
2.4.4 Limitations of the original method	4
3 Methods	5
3.1 Problem Formalization	5
3.2 Non Convex envelope for classification	5
3.3 Score Transformation	5
3.4 Envelope Constriction	5
3.4.1 CSA Two-Stage Calibration	5
3.4.2 τ Computation	6
3.4.3 Radial Method	6
3.4.4 Strip Method	6
3.4.5 1D	6
3.4.6 Class-Specific $\hat{t}$ Calculation	6
3.5 $\hat{t}$ Calculation	6
4 Data	7
4.1 Data Description	7
4.2 Protein to Phage Aggregation	7
4.2.1 NaN handling after aggregation	7
4.3 Train-Test Split	8
5 Experiments and Results	9
5.1 Synthetic Data Validation	9
5.1.1 Purpose and setup	9

5.1.2	2D Experiments	9
5.1.3	Extension to K Dimensions	9
5.1.4	Multiple Hosts	9
5.2	Gram Type Classification	10
5.2.1	Experimental Setup	10
5.2.2	Results	10
5.3	Host Classification	10
6	Conclusion and Future Work	11
6.1	Summary.....	11
6.2	Limitations	11

1 Introduction

The rapid rise of antibiotic-resistant bacteria is one of the major challenges that global medicine is facing today. To overcome this issue, researchers are going to alternative solutions, one of which is phage therapy. Phage therapy relies on using the naturally occurring viruses to target and destroy the harmful bacteria.

However, To make this therapy a viable alternative, we first need to understand the basic biology of these viruses. Bacteriophages, or phages are viruses that infect bacterial hosts that and are one of the most abundant entities on the planet. Each phage contains a collection of structural and functional proteins that collectively determine the bacterial strains that the phage can infect - this property is called the host range of the phage.

These phages are highly host-specific: a phage which infects one bacterial host may have very little to no effect on another. This particular property of theirs has naturally generated interest in modelling and applications in the context of..... . Therefore a significant problem in the phage biology is the prediction of the bacterial hosts associated with a given phage. Accurate host-predictions can have applications in

The host specificity of phages is one of phage therapy's greatest advantage, but at the same time, it also brings a very significant practical challenge. While it allows us to target specific harmful bacteria without destroying the beneficial ones, it also causes problems in finding exactly the appropriate phage for a specific infection. Doing this by only experimental screening can be time-consuming and expensive. This motivates the development of reliable computational models that are capable of predicting these phage-host interactions directly from the protein level information.

The existing computational tools, however have a very important limitation: they typically only produce point predictions, which are often represented as a scalar confidence score between 0 and 1 with no associated measure of uncertainty. This lack of having an uncertainty quantification is problematic in clinical contexts, where failing to identify a true host could mean missing a life-saving treatment. Consequently, to build trust in automated host prediction, we need a mathematical framework that is capable of quantifying its own predictive uncertainty and provides reliable confidence guarantees.

The primary objective of this report is to address this limitation by developing an uncertainty-aware framework for phage-host prediction using **Conformal Prediction**. Conformal prediction provides a distribution-free framework that allows us to transform standard point predictions into prediction sets that are guaranteed to contain the true bacterial host with a user-defined confidence level.

In this project, we build our phage representations by aggregating collections of protein-level functional scores. This process naturally produced high-dimensional, structured score vectors, as different proteins and functions vary in their data availability. These score vectors also often contain missing or noisy information. To handle these chal-

complete
sentence
(treat-
ment???)

find ap-
plica-
tions

lenges, we explored geometric conformal prediction methods that can construct reliable prediction sets even when dealing with incomplete biological data. As an intermediate step to validate our geometric approach in a controlled environment, we first evaluated the framework on a binary Gram-type classification task before moving on to the full multi-host prediction problem.

2 Background & State of the Art

2.1 Bacteriophage Host Prediction Problem

- What bacteriophages are: viruses that infect bacteria, most abundant biological entities on the planet
- The host range property: each phage encodes structural proteins that determine which bacterial strains it can infect
- Why host specificity is both an advantage and a limitation for phage therapy
- The computational problem: given a phage genome and a set of candidate hosts $\mathcal{H} = \{h_1, \dots, h_m\}$, identify the subset $\subseteq \mathcal{H}$ of true hosts
- Scale: hundreds of thousands of phage genomes, thousands of candidate hosts, experimental screening infeasible

2.2 State of the Art

already existing method for the phage host prediction

- Brief overview of existing approaches in machine learning
- Why Bayesian methods are unsuitable here: require distributional assumptions, computationally expensive for large-scale genomic data, posteriors depend heavily on prior choice
- Why calibration methods are insufficient: they produce calibrated probabilities but no formal finite-sample guarantee
- Why conformal prediction is the right choice: distribution-free, finite-sample guarantee, no assumptions beyond exchangeability, computationally lightweight

2.3 Conformal Prediction

2.3.1 Setup and nonconformity scores

- define exchangeability, define $s(x, y)$ -> explain the intuition

2.3.2 prediction set and coverage guarantee

- define $C(x)$

2.3.3 Split conformal prediction

- why we use it, the train/calibration/test separation, validity still holds

2.4 The CSA Method

2.4.1 Motivation

- from scalar to multivariate nonconformity scores -> why a scalar threshold is insufficient when the predictor produces a vector

2.4.2 Two-stage calibration

S_1 for shape discovery, S_2 for size scaling -> what each stage does and why they must be disjoint

2.4.3 The envelope concept

- a region $E \subset \mathbb{R}^K$ learned from S_1 , scaled by $\hat{t}$ calibrated on S_2

2.4.4 Limitations of the original method

- iterative search for $\hat{t}$, convexity assumptions, no class-conditional guarantee -> each of these is addressed in Chapter 3

3 Methods

3.1 Problem Formalization

- Input: a phage represented by a score vector in $[0, 1]^K$ where K is the number of functional dimensions
- Define the label space $\mathcal{Y}$ (finite set of classes)
- Goal: construct a prediction set $C(x) \subseteq Y$ with the conformal guarantee
- The one-envelope design: build one single envelope on all training phages under their true-label NC transformation, rather than separate envelopes per class

3.2 Non Convex envelope for classification

- In the regression setting of the original CSA paper, envelopes are typically convex because the score space is unbounded and the data does not cluster into separated groups
- In our classification setting, the score space is bounded in $[0, 1]^K$ and — crucially — true NC vectors for different classes occupy distinct, well-separated regions of this space (example)
- In the classification, produce tighter prediction sets without sacrificing the coverage guarantee
- The guarantee still holds because it depends only on the exchangeability of the data and the quantile construction, not on the shape of the envelope

3.3 Score Transformation

- General formulation: for each class $y \in \mathcal{Y}$, define a transformation $\phi_y: [0, 1]^K \rightarrow [0, 1]^K$ that maps the raw score vector to an conformal score vector such that a true infector of class y is near the origin
- Property we require: $\phi_{i_y}(\mathbf{s})$ is small when $\mathbf{s}$ is consistent with label y , large otherwise
- we make a single unified envelope coherent across all classes

3.4 Envelope Constriction

3.4.1 CSA Two-Stage Calibration

- Training vectors are split into two disjoint halves S_1 and S_2 (50/50 random split).
- S_1 (shape discovery): used to learn the geometry of the envelope.
- S_2 (size scaling): used to calibrate the scaling factor $\hat{t}$ that determines the final size of the envelope.

- The two stages must use disjoint data to preserve the exchangeability argument underlying the coverage guarantee.

3.4.2 τ Computation

- Define $\tau(\mathbf{s}) = \inf\{t > 0 : \mathbf{s} \in t \cdot E\}$, the minimum scaling factor needed to include point $\mathbf{s}$ in the scaled envelope E .
- The global scaling factor is then determined by the $(1 - \alpha)$ quantile of the τ values over S_2 :

$$\hat{t} = \tau_{(k)} \quad \text{where} \quad k = \lceil (|S_2| + 1)(1 - \alpha) \rceil$$

where $\tau_{(k)}$ denotes the k -th smallest value of τ in S_2 .

- Key contribution: for each of the three envelope geometries, $\tau(\mathbf{s})$ has a closed-form expression computable in a single pass — no iterative search required.
- Computational benefit: the entire calibration step reduces to one forward pass over S_2 followed by a sort.
- Contrast with the original CSA paper which uses iterative binary search.

3.4.3 Radial Method

3.4.4 Strip Method

3.4.5 1D

3.4.6 Class-Specific $\hat{t}$ Calculation

3.4.1 — CSA Two-Stage Calibration

3.5 $\hat{t}$ Calculation

- class conditional (why and how)

4 Data

4.1 Data Description

- The two files: protein-level score matrix (1,853,074 proteins $\times$ 40 functional score columns, values in [0,1], NaN when protein not annotated with that function) and meta-data file (same 1,853,074 rows, columns: accession, host, host_type, split)
- Scale of the dataset: 18,474 unique phages after aggregation
- The split column: integer 0–4 encoding a genome-cluster-based partition — phages in the same genome cluster are always in the same fold, ensuring test phages are genuinely different from training phages at the sequence level
- Class distribution: gram-neg vs gram-pos vs unknown
- The NaN structure: NaN is structural missingness, not noise — a protein that was not annotated with a given function simply has no score for that function, which is different from a score of zero

4.2 Protein to Phage Aggregation

- The raw data is at the protein level but conformal prediction operates at the phage level — aggregation is needed
- NaN-aware mean pooling: for each phage and each functional score column, compute the mean across all proteins of that phage that have a non-NaN value for that column. Proteins with NaN for a given column are excluded from the mean for that column only
- Why not replace NaN with zero before pooling: treating an unannotated protein as having zero evidence would push the phage’s aggregated score artificially toward zero, misrepresenting its biology
- After aggregation: result is a matrix of shape (number of phages) $\times$ 40, still with NaNs in cells where none of the phage’s proteins were annotated with that function

4.2.1 NaN handling after aggregation

- After mean pooling, some phages still have NaN in certain columns -> this happens when none of their proteins were annotated with a particular function
- Strategy: fill remaining NaNs using column-wise medians computed strictly on the training set (never the test set, to prevent data leakage) - The median cap: any training median above 0.6 is capped at 0.5 -> this prevents an imputed value from falsely pushing a phage toward a particular class - any remaining NaNs after median imputation are filled with 0.5 (neutral value)

4.3 Train-Test Split

- Standard split: split 0 is the held-out test set, splits 1–4 form the training set
- The contamination problem: explain the protein cluster (pc) column -> a protein's cluster may include proteins from phages in different folds. The models that produced the scores were trained on protein clusters, so a protein from a training-fold phage may share a cluster with a test-fold protein, introducing subtle leakage
- The masking fix: the file `all_random_pc_splits.pkl` encodes for each protein and each (host, function) model which fold that protein was assigned to. Setting scores to NaN where the protein was not in the test set of the relevant model, before aggregation, removes the contaminated predictions. The existing NaN-aware mean pooling then handles these naturally -> this is the "missing value" framing

5 Experiments and Results

5.1 Synthetic Data Validation

5.1.1 Purpose and setup

- Why synthetic validation before real data: allows controlled verification that the coverage guarantee holds exactly, with known ground truth distributions
- General setup: simulate score vectors from known distributions, apply all three envelope methods, verify empirical coverage matches theoretical target $1 - \alpha$

5.1.2 2D Experiments

- Three scenarios tested in 2D ($K=2$):
 - Uncorrelated scores: two independent uniform or Gaussian dimensions
 - Correlated scores: scores with positive correlation -> tests whether the envelope adapts to the correlation structure
 - Asymmetric predictors: one informative dimension and one uninformative dimension -> tests whether the envelope correctly ignores the noisy dimension
- For each scenario: show the envelope shape visually, report empirical coverage across multiple runs, confirm it matches $1 - \alpha$
- Key observation from 2D: the radial and strip envelopes are visibly non-convex and wrap tightly around the data, while collapsed 1D produces a simple interval

5.1.3 Extension to K Dimensions

- Repeat the experiments for $K \in 2, 5, 10, 20$ with a single host
- Show that all three methods maintain coverage $\geq 1 - \alpha$ across all dimensionalities
- Discuss the effect of increasing K on envelope quality: radial method becomes sparser (fewer S1 points per sector), strip method has more dimension pairs to condition on, collapsed 1D is unaffected
- Show that the analytical τ computation gives identical results to what iterative search would give, confirming correctness

5.1.4 Multiple Hosts

- Extend to a simulated multi-class setting with multiple hosts ($|\mathcal{Y}| = 6, K = 5$)
- Each host's true score distribution is a distinct cluster in $[0, 1]^K$
- Verify that the single unified envelope correctly assigns the true host to the prediction set with probability $\geq 1 - \alpha$
- Show 3D visualisation of the envelope for $K=3$

5.2 Gram Type Classification

5.2.1 Experimental Setup

- Data: split-0 phages only (80% calibration, 20% test), unknown host types excluded
- Parameters: $\alpha = 0.1$, $M = 200$ directions (radial), $M=20$ bins (strip), 10 independent runs to estimate variance
- The NC transformation used: keep scores as-is for one class, flip with $1 - s$ for the other

5.2.2 Results

- table: overall coverage, per-class coverage, average set size, empty set rate, size-2 rate, singleton accuracy std across 10 runs for all three methods
- Bar charts: set size distribution, overall and per-class coverage

5.3 Host Classification

6 Conclusion and Future Work

6.1 Summary.....

6.2 Limitations

- gram-pos coverage -